

第九讲：深度强化学习



主讲人 陈达贵

清华大学自动化系
在读硕士



深蓝学院
shenlanxueyuon.com

强化学习理论与实践



目录


1 本章简介

2 Value-based DRL

3 Policy-based DRL

4 Trust Region based DRL

章节目录




本章目录

- 1 本章简介
- 2 Value-based DRL
- 3 Policy-based DRL
- 4 Trust Region based DRL

深度强化学习简介

- 深度学习 + 强化学习 = 深度强化学习
- 使用深度神经网络作为强化学习的函数近似器
- 深度学习本质上是一种表示学习，它负责将原始数据表达为和问题相关的特征
- 例：flappybird
 - 图像输入：DRL
 - 坐标输入：RL



讲解流程

- 基于值函数: 从 DQN 到 Rainbow
- 基于策略
 - 策略梯度框架
 - 信赖域优化

本节课会大量引用文献，为了方便大家得文献阅读，尽量使用和原文类似的英文表达



目录

1 本章简介

2 Value-based DRL

3 Policy-based DRL

4 Trust Region based DRL


DQN

- Playing Atari with Deep Reinforcement Learning (NIPS2013)
- Human-level control through deep reinforcement learning (Nature2015)
- 重要意义
 - 不用函数近似无法解决大规模问题，用函数近似训练不稳定
 - 首次证明了能够通过 raw pixels 解决游戏问题
 - 对所有游戏通用

DQN

■ 关键特点

- Q Learning+DNN
- Experience Replay
- Target Network



DQN

Algorithm 1: deep Q-learning with experience replay.Initialize replay memory D to capacity N Initialize action-value function Q with random weights θ Initialize target action-value function \hat{Q} with weights $\theta^- = \theta$ **For** episode = 1, M **do**Initialize sequence $s_1 = \{x_1\}$ and preprocessed sequence $\phi_1 = \phi(s_1)$ **For** $t = 1, T$ **do**With probability ε select a random action a_t otherwise select $a_t = \operatorname{argmax}_a Q(\phi(s_t), a; \theta)$ Execute action a_t in emulator and observe reward r_t and image x_{t+1} Set $s_{t+1} = s_t, a_t, x_{t+1}$ and preprocess $\phi_{t+1} = \phi(s_{t+1})$ Store transition $(\phi_t, a_t, r_t, \phi_{t+1})$ in D Sample random minibatch of transitions $(\phi_j, a_j, r_j, \phi_{j+1})$ from D

$$\text{Set } y_j = \begin{cases} r_j & \text{if episode terminates at step } j+1 \\ r_j + \gamma \max_{a'} \hat{Q}(\phi_{j+1}, a'; \theta^-) & \text{otherwise} \end{cases}$$
Perform a gradient descent step on $(y_j - Q(\phi_j, a_j; \theta))^2$ with respect to the network parameters θ Every C steps reset $\hat{Q} = Q$ **End For****End For**<http://blog.csdn.net/yeqiang19910412>

Double DQN


- Deep Reinforcement Learning with Double Q-learning (AAAI2016)
- Q 学习中存在过估计
 - Q 学习中的 TD 目标值 $r + \gamma \max_a Q(s', a)$ 中存在 \max 操作
 - 这会引入一个正向的偏差
 - 因此建模两个 Q 网络，一个用于选动作，一个用于评估动作

$$r + \gamma Q^B(s', \arg \max_a Q^A(s', a))$$

例子

对于状态 s 下，如果对于所有的 a ，真实的 $q(s, a)$ 均为 0，但是估计值由于不精确，会导致有些大于 0，有些小于 0。对估计值取最大，会引入一个正向的偏差

Double DQN



Algorithm 1 Double Q-learning

```

1: Initialize  $Q^A, Q^B, s$ 
2: repeat
3:   Choose  $a$ , based on  $Q^A(s, \cdot)$  and  $Q^B(s, \cdot)$ , observe  $r, s'$ 
4:   Choose (e.g. random) either UPDATE(A) or UPDATE(B)
5:   if UPDATE(A) then
6:     Define  $a^* = \arg \max_a Q^A(s', a)$ 
7:      $Q^A(s, a) \leftarrow Q^A(s, a) + \alpha(s, a) (r + \gamma Q^B(s', a^*) - Q^A(s, a))$ 
8:   else if UPDATE(B) then
9:     Define  $b^* = \arg \max_a Q^B(s', a)$ 
10:     $Q^B(s, a) \leftarrow Q^B(s, a) + \alpha(s, a) (r + \gamma Q^A(s', b^*) - Q^B(s, a))$ 
11:   end if
12:    $s \leftarrow s'$ 
13: until end
  
```

Dueling DQN

- Dueling Network Architectures for Deep Reinforcement (ICML2016)
- 将 Q 函数分解成 V 函数和优势 (A) 函数





Figure 1. A popular single stream Q -network (**top**) and the dueling Q -network (**bottom**). The dueling network has two streams to separately estimate (scalar) state-value and the advantages for each action; the green output module implements equation (9) to combine them. Both networks output Q -values for each action.

Dueling DQN

■ 优势

- 对于很多状态并不需要估计每个动作的值，增加了 V 函数的学习机会
- V 函数的泛化性能好，当有新动作加入时，并不需要重新学习
- 减少了 Q 函数由于状态和动作维度差导致的噪声和突变



Prioritized Experience Replay

- Prioritized Experience Replay (ICLR 2016)
- DQN 算法的一个重要改进是 Experience Replay
- 训练时从 Memory 中均匀采样
- Prioritized Experience Replay 就是维护了一个带优先级的 Experience Replay
 - 不同的 Experience 的权重不同
 - 用 TD 误差去衡量权重
 - 需要使用 sum-tree 以及 binary heap data structure 去实现
 - 新的 transition 的 TD 误差会被设置为最大
 - 类似于 DP 中的优先清理
- Experience Replay 使得更新不受限于实际经验的顺序, Prioritized Experience Replay 使得更新不受限于实际经验的频率

Prioritized Experience Replay


具体使用时会存在一些问题

- TD 误差对噪声敏感
- TD 误差小的 transition 长时间不更新
- 过分关注 TD 误差大的 transition 丧失了样本多样性
- 使用某种分布采样了 Experience, 会引入 Bias

解决方法

- 两种变体: $p_i = |\delta_i| + \varepsilon$ 或者 $p_i = \frac{1}{rank(i)}$
- 加入重要性采样来消除 Bias

Prioritized Experience Replay

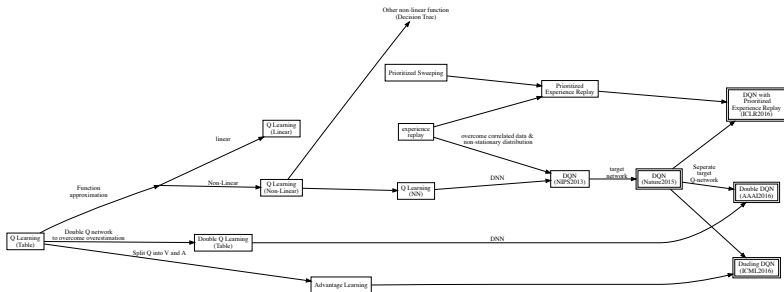


Prioritized Experience Replay

Algorithm 1 Double DQN with proportional prioritization

```
1: Input: minibatch  $k$ , step-size  $\eta$ , replay period  $K$  and size  $N$ , exponents  $\alpha$  and  $\beta$ , budget  $T$ .
2: Initialize replay memory  $\mathcal{H} = \emptyset$ ,  $\Delta = 0$ ,  $p_1 = 1$ 
3: Observe  $S_0$  and choose  $A_0 \sim \pi_\theta(S_0)$ 
4: for  $t = 1$  to  $T$  do
5:   Observe  $S_t, R_t, \gamma_t$ 
6:   Store transition  $(S_{t-1}, A_{t-1}, R_t, \gamma_t, S_t)$  in  $\mathcal{H}$  with maximal priority  $p_t = \max_{i < t} p_i$ 
7:   if  $t \equiv 0 \pmod K$  then
8:     for  $j = 1$  to  $k$  do
9:       Sample transition  $j \sim P(j) = p_j^\alpha / \sum_i p_i^\alpha$ 
10:      Compute importance-sampling weight  $w_j = (N \cdot P(j))^{-\beta} / \max_i w_i$ 
11:      Compute TD-error  $\delta_j = R_j + \gamma_j Q_{\text{target}}(S_j, \arg \max_a Q(S_j, a)) - Q(S_{j-1}, A_{j-1})$ 
12:      Update transition priority  $p_j \leftarrow |\delta_j|$ 
13:      Accumulate weight-change  $\Delta \leftarrow \Delta + w_j \cdot \delta_j \cdot \nabla_\theta Q(S_{j-1}, A_{j-1})$ 
14:    end for
15:    Update weights  $\theta \leftarrow \theta + \eta \cdot \Delta$ , reset  $\Delta = 0$ 
16:    From time to time copy weights into target network  $\theta_{\text{target}} \leftarrow \theta$ 
17:  end if
18:  Choose action  $A_t \sim \pi_\theta(S_t)$ 
19: end for
```


小结



Rainbow

- Rainbow: Combining Improvements in Deep Reinforcement Learning
- 集合了多种 DQN 的变体
 - DQN
 - Double DQN
 - Dueling DQN
 - Prioritized Experience Replay
 - NoiseNet: (Noisy Networks for Exploration, AAAI2018)
 - Distributional RL: (A Distributional Perspective on Reinforcement Learning, ICML2017)

Rainbow





目录

- 1 本章简介
- 2 Value-based DRL
- 3 Policy-based DRL**
- 4 Trust Region based DRL

DPG

- Deterministic Policy Gradient Algorithms (ICML2014)
- 将策略梯度定理用到了高维和连续动作
- 常规的策略梯度方法无法用到高维和连续动作空间
 - 连续动作无法使用概率分布输出的形式
 - 高维动作空间中采样费时
 - 随机策略梯度是通过采样的方式估算策略梯度, 需要在状态空间和动作空间内采样

$$\nabla_{\pi_{\theta}} J(\theta) = \int_S \int_{\mathcal{A}} \rho^{\pi_{\theta}}(s) \nabla_{\theta} \log \pi_{\theta}(a|s) Q^{\pi_{\theta}}(s, a) ds da$$

- 直接采用确定性策略输出

$$a = \pi(s)$$

- 由此诞生几个问题:
 - 确定性策略情况下如何求策略梯度?
 - 如何探索?

DPG


- 过去一直认为无模型情况下确定性策略梯度不存在
- DPG 证明了确定性策略梯度定理的存在, 建立了它跟 Q 函数梯度的关系

$$\nabla_{\theta} J(\theta) = \int_S \rho^{\pi_{\theta}}(s) \nabla_{\theta} \pi_{\theta} \nabla_a Q^{\pi_{\theta}}(s, a)|_{a=\pi_{\theta}(s)} ds$$

- 只需要对状态 S 进行积分
- 使用 Off-policy 的方式探索并更新 (需要重要性采样)

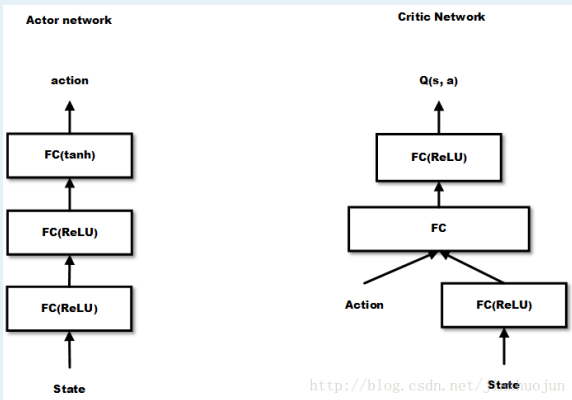
DDPG

- Continuous Control with Deep Reinforcement Learning (ICRL2016)
- 结合了 DQN 和 DPG
- 使用了 DQN 的两种技术：Experience Replay 和 Target Network
- 利用随机过程产生探索性动作



DDPG

建模方式



DDPG

Algorithm 1 DDPG algorithm

Randomly initialize critic network $Q(s, a|\theta^Q)$ and actor $\mu(s|\theta^\mu)$ with weights θ^Q and θ^μ .

Initialize target network Q' and μ' with weights $\theta^{Q'} \leftarrow \theta^Q, \theta^{\mu'} \leftarrow \theta^\mu$

Initialize replay buffer R

for episode = 1, M **do**

 Initialize a random process \mathcal{N} for action exploration

 Receive initial observation state s_1

for t = 1, T **do**

 Select action $a_t = \mu(s_t|\theta^\mu) + \mathcal{N}_t$ according to the current policy and exploration noise

 Execute action a_t and observe reward r_t and observe new state s_{t+1}

 Store transition (s_t, a_t, r_t, s_{t+1}) in R

 Sample a random minibatch of N transitions (s_i, a_i, r_i, s_{i+1}) from R

 Set $y_i = r_i + \gamma Q'(s_{i+1}, \mu'(s_{i+1}|\theta^{\mu'}))|\theta^{Q'}$

 Update critic by minimizing the loss: $L = \frac{1}{N} \sum_i (y_i - Q(s_i, a_i|\theta^Q))^2$

 Update the actor policy using the sampled policy gradient:

$$\nabla_{\theta^\mu} J \approx \frac{1}{N} \sum_i \nabla_a Q(s, a|\theta^Q)|_{s=s_i, a=\mu(s_i)} \nabla_{\theta^\mu} \mu(s|\theta^\mu)|_{s_i}$$

 Update the target networks:

$$\theta^{Q'} \leftarrow \tau \theta^Q + (1 - \tau) \theta^{Q'}$$

$$\theta^{\mu'} \leftarrow \tau \theta^\mu + (1 - \tau) \theta^{\mu'}$$


end for

end for

<http://blog.csdn.net/shenshikexmu>

A3C

- Asynchronous Methods for Deep Reinforcement Learning (ICML2016)
- Online 的算法和 DNN 结合后不稳定 (样本关联性)
- 通过创建多个 agent 在多个环境执行异步学习构建 batch
 - 来自不同环境的样本无相关性
 - 不依赖于 GPU 和大型分布式系统
 - 不同线程使用了不同的探索策略, 增加了探索量



A3C


Algorithm S3 Asynchronous advantage actor-critic - pseudocode for each actor-learner thread.*// Assume global shared parameter vectors θ and θ_v and global shared counter $T = 0$* *// Assume thread-specific parameter vectors θ' and θ'_v* Initialize thread step counter $t \leftarrow 1$ **repeat**Reset gradients: $d\theta \leftarrow 0$ and $d\theta_v \leftarrow 0$.Synchronize thread-specific parameters $\theta' = \theta$ and $\theta'_v = \theta_v$ $t_{start} = t$ Get state s_t **repeat**Perform a_t according to policy $\pi(a_t|s_t; \theta')$ Receive reward r_t and new state s_{t+1} $t \leftarrow t + 1$ $T \leftarrow T + 1$ **until** terminal s_t **or** $t - t_{start} == t_{max}$ $R = \begin{cases} 0 & \text{for terminal } s_t \\ V(s_t, \theta'_v) & \text{for non-terminal } s_t // \text{ Bootstrap from last state} \end{cases}$ **for** $i \in \{t-1, \dots, t_{start}\}$ **do** $R \leftarrow r_i + \gamma R$ Accumulate gradients wrt θ' : $d\theta \leftarrow d\theta + \nabla_{\theta'} \log \pi(a_i|s_i; \theta')(R - V(s_i; \theta'_v))$ Accumulate gradients wrt θ'_v : $d\theta_v \leftarrow d\theta_v + \partial (R - V(s_i; \theta'_v))^2 / \partial \theta'_v$ **end for**Perform asynchronous update of θ using $d\theta$ and of θ_v using $d\theta_v$.**until** $T > T_{max}$

A2C

- Openai Blog(<https://blog.openai.com/baselines-acktr-a2c/>)
- 改异步为同步, 能够更好地利用 GPU
- 在 batch_size 较大时效果好



小结





目录

- 1 本章简介
- 2 Value-based DRL
- 3 Policy-based DRL
- 4 Trust Region based DRL

简要介绍

策略梯度算法的更新步长很重要

- 步长太小，导致更新效率低下
- 步长太大，导致参数更新的策略比上次更差，通过更差的策略采样得到的样本更差，导致学习再次更新的参数会更差，最终崩溃

如何选择一个合适的步长，使得每次更新得到的新策略所实现的回报值单调不减

- 信赖域 (Trust Region) 方法指在该区域内更新，策略所实现的回报值单调不减

一些基础

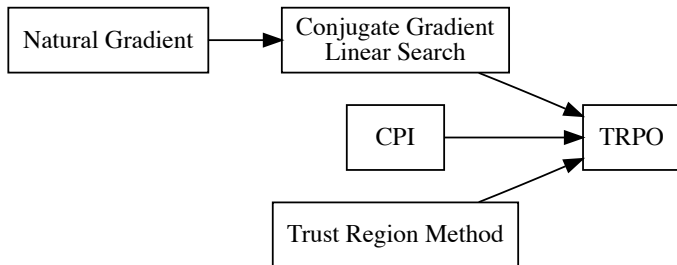
- 自然梯度: Natural Gradient Works Efficiently in Learning, 1998
 - 在黎曼空间里面, 最快的下降方向不是梯度方向, 而是自然梯度方向 $G^{-1}(\theta)J(\theta)$
 - 只有当坐标系统正交, 才退化成欧式空间
 - 神经网络中的参数空间是黎曼空间
 - 其中 G 为 Reimannian metric tensor
 - 统计问题中, G 可以用 Hessian 矩阵去计算

一些基础

- 保守策略迭代, CPI: Approximately Optimal Approximate Reinforcement Learning, 2002
 - 给出策略性能增长的条件
 - 策略更新后的所有优势函数非负
 - 使用混合更新的方式更新策略


TRPO

- Trust Region Policy Optimization, ICML2015
- 以 CPI 为基础，推导出策略更新后性能的下界，通过优化下界优化原函数
- 实际操作时用 KL 散度作为约束
- 求解带约束的优化问题时，利用自然梯度
- 实际求解是利用了共轭梯度 + 线性搜索的方法，避免求自然梯度



PPO

- Proximal Policy Optimization Algorithms, 2017
- Openai blog(<https://blog.openai.com/openai-baselines-ppo/>)
- TRPO 太复杂，普通 PG 效果又不好
- PPO 本质上是 TRPO 的简化版
- 移除了 KL 惩罚项和交替更新
- 由于性能好，且容易实现，已经成为默认的 OPENAI 算法



PPO

$$L^{CLIP}(\theta) = \hat{E}_t [\min(r_t(\theta)\hat{A}_t, \text{clip}(r_t(\theta), 1 - \varepsilon, 1 + \varepsilon)\hat{A}_t)]$$

- θ is the policy parameter
- \hat{E}_t denotes the empirical expectation over timesteps
- r_t is the ratio of the probability under the new and old policies, respectively
- \hat{A}_t is the estimated advantage at time t
- ε is a hyperparameter, usually 0.1 or 0.2

其他信赖域算法

- ACKTR: Scalable trust-region method for deep reinforcement learning using Kronecker-factored approximation
- ACER: Sample Efficient Actor-Critic with Experience Replay
- GAE: High-Dimensional Continuous Control Using Generalized Advantage Estimation
- ...